

## Swift 5.5 20 Sept 2021

- Async/ await  
SE-0296
- Async let  
SE-0317
- Async Sequences  
SE-0298
- Structured concurrency  
SE-0304
- Actors  
SE-0306
- Global Actors  
SE-0316
- Async/ await / async let interoperability with ObjC  
SE-0297
- Interfacing async with callbacks  
SE-0300
- AsyncStreams  
SE-0314
- Tasks local values  
SE-0311
- Actor isolation control  
SE-0313

## Swift 5.5.2 Dec 14, 2021

Clarification of the semantics of Task where the main function runs  
SE-0323

## Swift 5.6 Mar 14, 2022

Remove sendable conformance in unsafe pointers  
SE-0331

Support incremental migration of concurrency checks  
SE-0337

## Swift 5.7 Sep 12, 2022

Sendable + closures  
SE-0302      Distributed Actors  
SE-0336

Clarify how non-isolated async Actor functions should work  
SE-0338

NoAsync - mark async unavailable  
SE-0340

Detailing how concurrency work on top level code  
SE-0343

Specifying how the runtime should work for distributed actors  
SE-0344

## Swift 5.9 Sep 18, 2023

Discarding task groups. Improving limitations of SE-304  
SE-0381

AsyncStream factory. Improving SE-0314  
SE-0388

Custom actor executor  
SE-0392

## Swift 5.10 Mar 5, 2024

Specifying how init and deinit should work on actors  
SE-0327

Isolation for default property values and default parameters values.  
SE-0411

Strict concurrent for global variables  
SE-0412

## Swift 6.0 Sep 17, 2024

Regions based isolation  
SE-0414

Task executor preferences  
SE-0417

Inferring sendable for keyPaths and methods  
SE-0418

Changes in concurrency rules for global actor isolated types  
SE-0434

AssumesIsolated / Dynamic isolation checks for non-strict concurrency code  
SE-0423

Custom isolation checks for SerialExecutor  
SE-0424

"sending" parameters and results values  
SE-0430

Giving async functions actor isolation / inheritance of actor isolation  
SE-0420

isolated(any) function type  
SE-0431

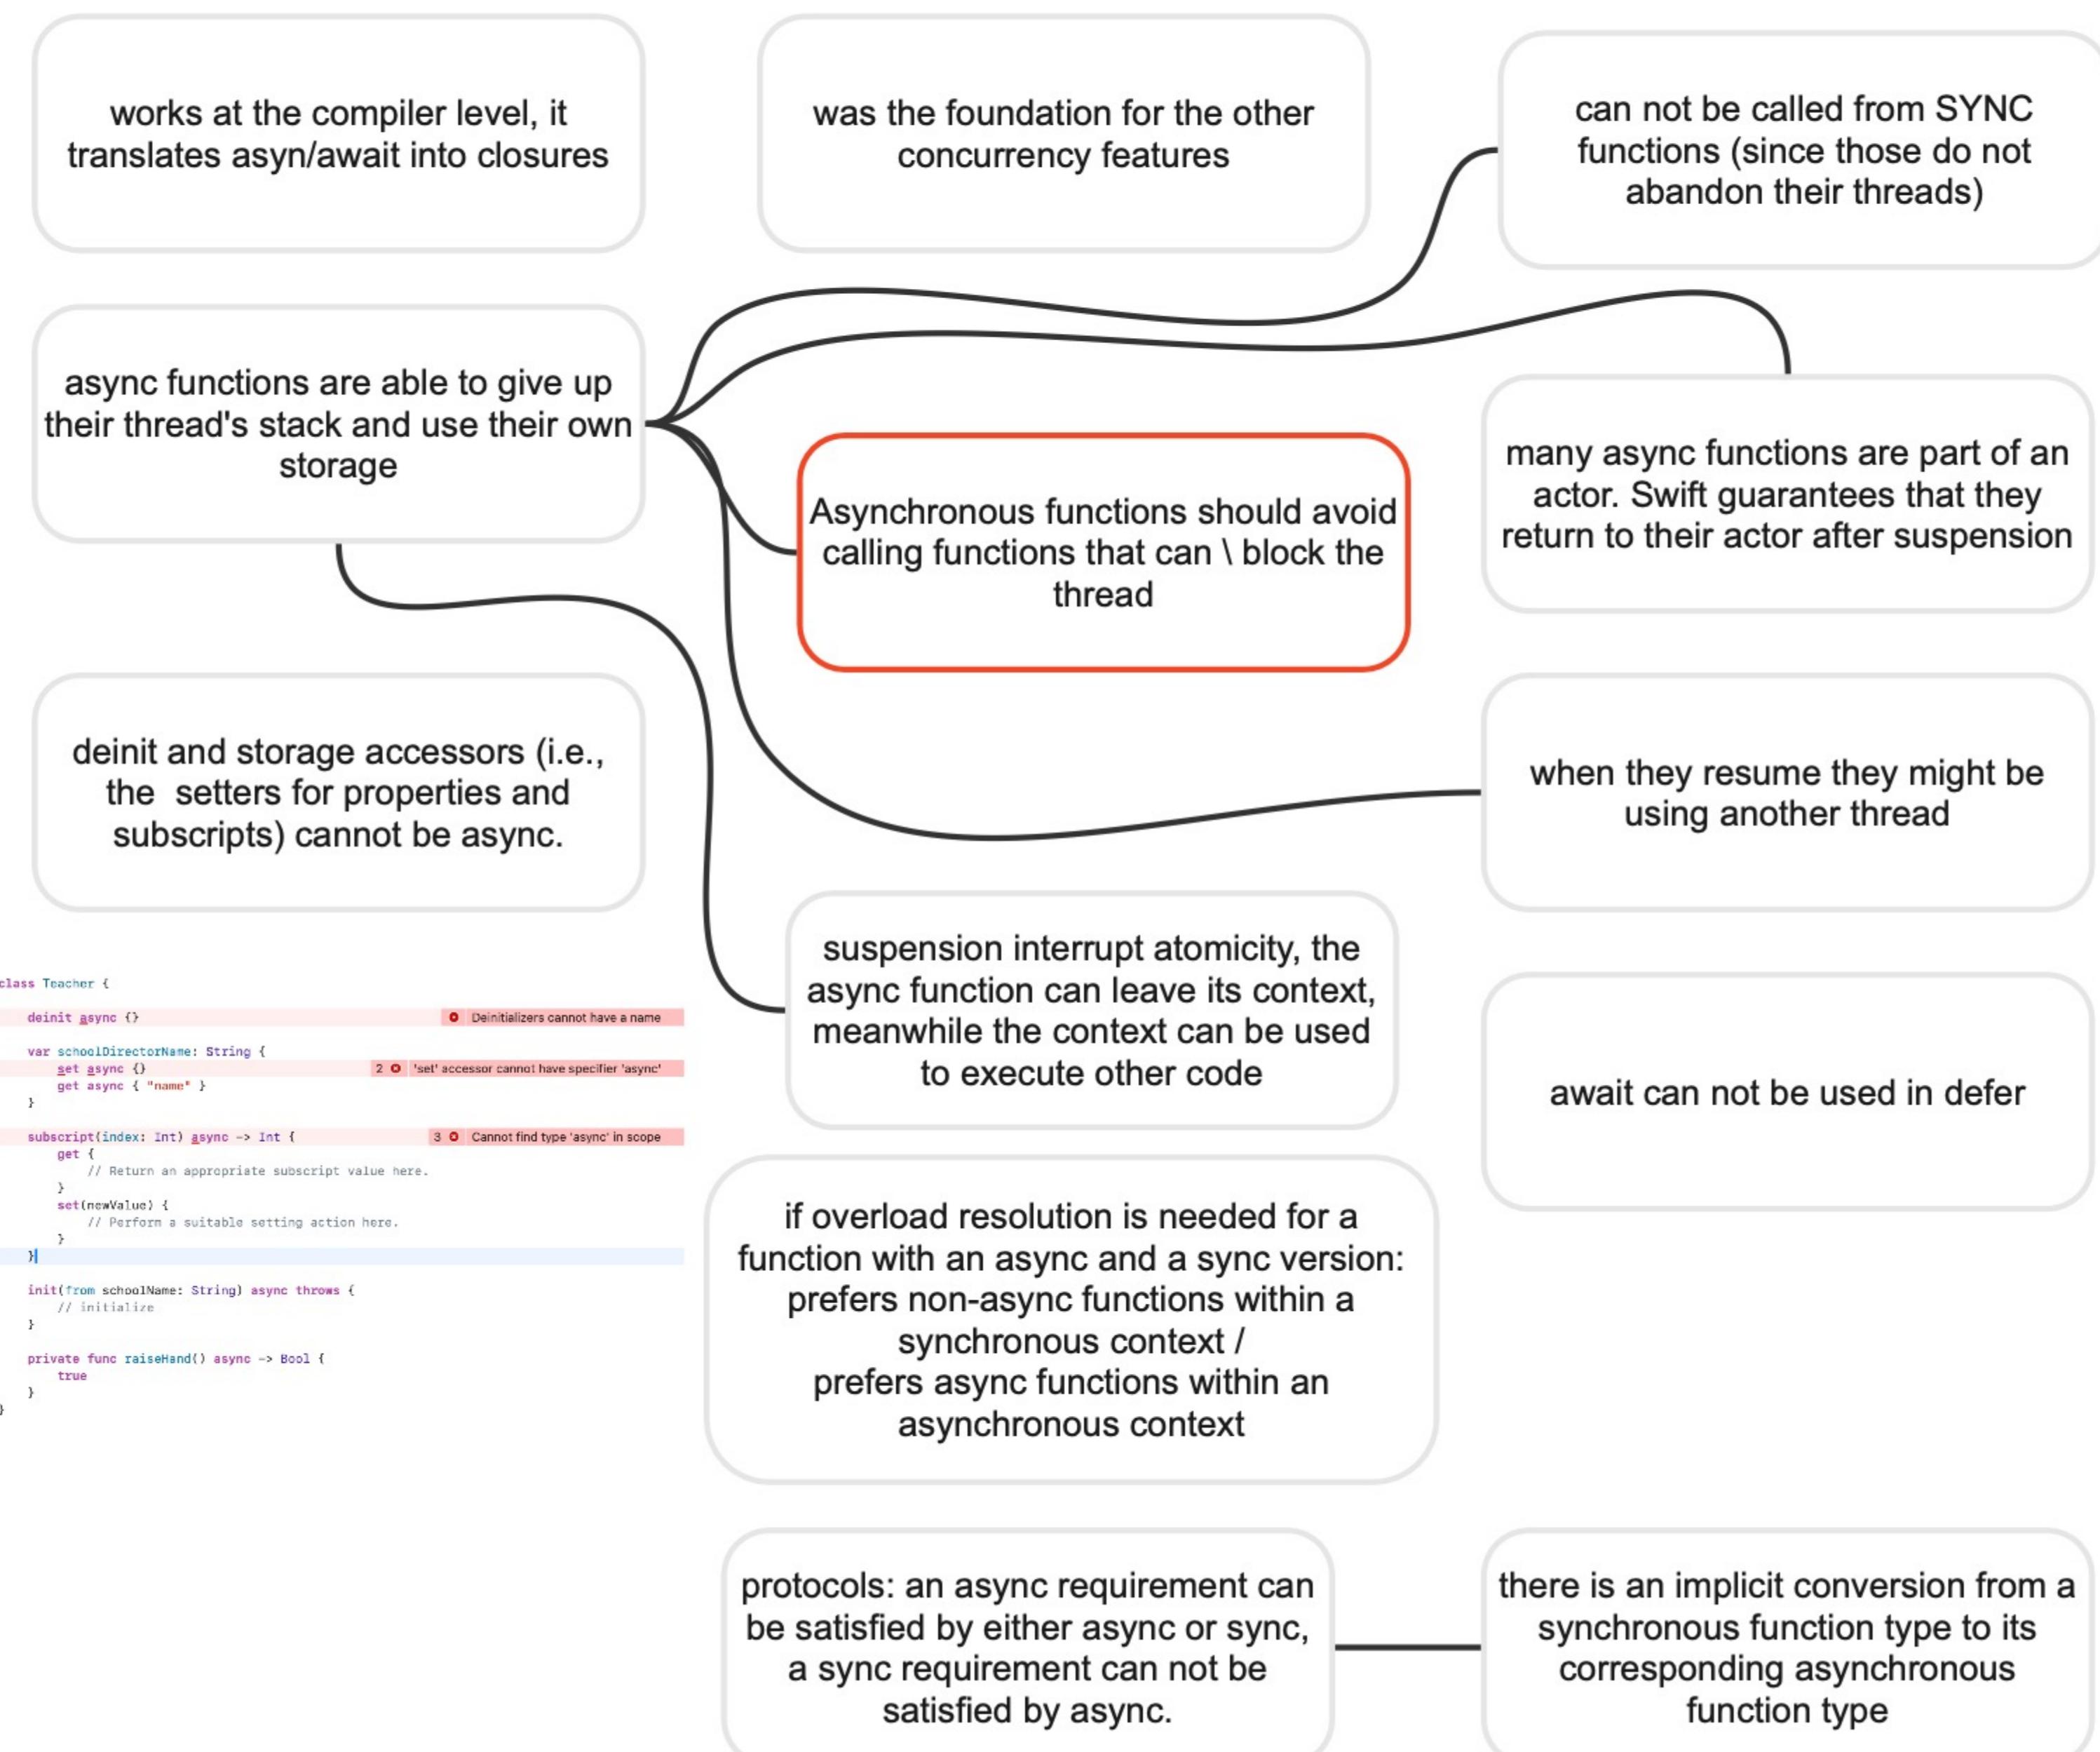
## Swift 6.1

Isolated sync deinit for actors and global actors. Lifting restriction from SE-0327  
SE-0371

Inferring TaskGroup child results  
SE-0442

Allow preventing global-actor inferences, using nonisolated  
SE-0449

# async await SE-0296



# objc interoperability SE-0297

Swift provides bridging from ObjC to  
async await

methods in ObjC are automatically  
converted to async when they fulfill  
some conditions (See proposal  
SE-0297)

functions in Swift can be marked with  
@objc and the compiler will translate  
them into ObjC callbacks, when  
possible

an actor can not inherit from a class,  
with the exception of NSObject, to  
precisely provide bridging.

An actor function can be marked  
@objc only if it is async or is non  
isolated. ObjectiveC does not know  
how to provide actor isolation of sync  
functions

markers were added to help the  
compiler and annotate old ObjC code  
with for example `swint_async` and  
`swift_async_error` (see SE-0297 for  
the whole list of attributes)

# async let SE-0317

was introduced as a way to hide the complexities of setting up child tasks and waiting for their results

the "let" indicates that it is a local constant the "async" part mean that it is evaluated in a concurrent child task

By default, child tasks use the **global, width-limited, concurrent executor**

the child task closure created is sendable and non-isolated

It is illegal to declare an async var

if you do not wait for an async let, the function still will take as much time to return as the longest of its child tasks takes to complete

Special attention needs to be given to the async let \_ = ...it will run and be cancelled and awaited-on implicitly, as the scope it was declared in is about to exit

It is *not* legal to escape a async let value to an escaping closure, because structures backing the async let implementation may be allocated on the stack rather than the heap. This makes them very efficient, however, make it unsafe to pass them to any escaping contexts

Limitation: the number of child-tasks is *dynamic*, it is not possible to express using async let because we'd have to know how many async let declarations to create at *compile time*

Limitation: Because async let declarations must be awaited on it is not possible to express "whichever completes first", and a task group must be used to implement such API

```
func makeDinner() async throws -> Meal {  
    async let veggies = chopVegetables()  
    async let meat = marinateMeat()  
    async let oven = preheatOven(temperature: 350)  
  
    let dish = Dish(ingredients: await [try veggies, meat])  
    return try await oven.cook(dish, duration: .hours(3))  
}
```

# closure + async await SE-0300

continuations are a way to mix  
callback based code  
with async code

swift provides continuation APIs,  
which are given to async functions.  
Inside closure based code can be  
called

when the closure based code  
finished the continuation must be  
resumed with success or error

resume must be called  
ONLY ONCE

the closure based code runs in the  
current task's context

if the continuation is unsafe, and the  
resume is called more than once,  
Swift produces undefined behaviour

it is better to use checked continuations, since it is a wrapper on top of unsafe, which checks how many times the resume is called, and also check for forgotten resumes, plus trap at runtime if there is any error in the handling of the continuation

```
func buyVegetables(  
    shoppingList: [String],  
    // a) if all veggies were in store, this is invoked *exactly-once*  
    onGotAllVegetables: ([Vegetable]) -> (),  
  
    // b) if not all veggies were in store, invoked one by one *one or more times*  
    onGotVegetable: (Vegetable) -> (),  
    // b) if at least one onGotVegetable was called *exactly-once*  
    // this is invoked once no more veggies will be emitted  
    onNoMoreVegetables: () -> (),  
  
    // c) if no veggies _at all_ were available, this is invoked *exactly once*  
    onNoVegetablesInStore: (Error) -> ()  
)  
// returns 1 or more vegetables or throws an error  
func buyVegetables(shoppingList: [String]) async throws -> [Vegetable] {  
    try await withUnsafeThrowingContinuation { continuation in  
        var veggies: [Vegetable] = []  
  
        buyVegetables(  
            shoppingList: shoppingList,  
            onGotAllVegetables: { veggies in continuation.resume(returning: veggies) },  
            onGotVegetable: { v in veggies.append(v) },  
            onNoMoreVegetables: { continuation.resume(returning: veggies) },  
            onNoVegetablesInStore: { error in continuation.resume(throwing: error) },  
        )  
    }  
}  
  
let veggies = try await buyVegetables(shoppingList: ["onion", "bell pepper"])
```

```
// returns 1 or more vegetables or throws an error  
func buyVegetables(shoppingList: [String]) async throws -> [Vegetable] {  
    try await withCheckedThrowingContinuation { continuation in  
        var veggies: [Vegetable] = []  
  
        buyVegetables(  
            shoppingList: shoppingList,  
            onGotAllVegetables: { veggies in continuation.resume(returning: veggies) },  
            onGotVegetable: { v in veggies.append(v) },  
            onNoMoreVegetables: { continuation.resume(returning: veggies) },  
            onNoVegetablesInStore: { error in continuation.resume(throwing: error) },  
        )  
    }  
}
```

# Async sequences SE-0298

async/await functions return 1 value...  
async/await sequences are the provided Swift way, to return many overtime

the main point is been able to return more than one value, along the time, without waiting for all the values to be calculated or read

the 2 protocols for sequences, provided by the standard lib are: AsyncSequence and AsyncIteratorProtocol

the compiler produces code to support using for...in with async sequences

The AsyncIterator MUST check for Task isCancelled and decide if to return the partially calculated result or if to return nil

After an AsyncIteratorProtocol types returns nil or throws an error from its next() method, all future calls to next() must return nil

```
public protocol AsyncSequence {
    associatedtype AsyncIterator: AsyncIteratorProtocol where AsyncIterator.Element == Element
    associatedtype Element
    __consuming func makeAsyncIterator() -> AsyncIterator
}

public protocol AsyncIteratorProtocol {
    associatedtype Element
    mutating func next() async throws -> Element?
}
```

async sequences come with some already defined cool functions

Function	Note
<code>contains(_ value: Element) async rethrows -&gt; Bool</code>	Requires Equatable element
<code>contains(where: (Element) async throws -&gt; Bool) async rethrows -&gt; Bool</code>	The <code>async</code> on the closure allows optional async behavior, but does not require it
<code>allSatisfy(_ predicate: (Element) async throws -&gt; Bool) async rethrows -&gt; Bool</code>	
<code>first(where: (Element) async throws -&gt; Bool) async rethrows -&gt; Element?</code>	
<code>min() async rethrows -&gt; Element?</code>	Requires Comparable element
<code>min(by: (Element, Element) async throws -&gt; Bool) async rethrows -&gt; Element?</code>	
<code>max() async rethrows -&gt; Element?</code>	Requires Comparable element
<code>max(by: (Element, Element) async throws -&gt; Bool) async rethrows -&gt; Element?</code>	
<code>reduce&lt;T&gt;(_ initialResult: T, _ nextPartialResult: (T, Element) async throws -&gt; T) async rethrows -&gt; T</code>	
<code>reduce&lt;T&gt;(into initialResult: T, _ updateAccumulatingResult: (inout T, Element) async throws -&gt; () ) async rethrows -&gt; T</code>	

# Async streams SE-0314

AsyncStream and AsyncThrowingStream were added to help starting to write more async sequences

Continuations were added to support connecting callback based code to async function. BUT this only worked for 1 single value returned... so something new was needed for closure based code which returns many values over time

when creating an stream, you receive a continuation, but this one can be called multiple time to YIELD elements

elements yielded are accumulated and buffered, until a consumer receives them

a continuation can also be finished. Then any accumulated value is given to the consumer before terminating with nil or throwing an exception

```
func buyVegetables(  
    shoppingList: [String],  
  
    // a) invoked once for each vegetable in the shopping list  
    onGotVegetable: (Vegetable) -> Void,  
  
    // b) invoked once all available veggies have been retrieved  
    onAllVegetablesFound: () -> Void,  
  
    // c) invoked if a non-vegetable food item is encountered  
    // in the shopping list  
    onNonVegetable: (Error) -> Void  
)  
  
// Returns a stream of veggies  
func findVegetables(shoppingList: [String]) -> AsyncThrowingStream<Vegetable> {  
    AsyncThrowingStream { continuation in  
        buyVegetables(  
            shoppingList: shoppingList,  
            onGotVegetable: { veggie in continuation.yield(veggie) },  
            onAllVegetablesFound: { continuation.finish() },  
            onNonVegetable: { error in continuation.finish(throwing: error) }  
        )  
    }  
}
```

# Async streams and View Controllers

When using AsyncStream, or a long running task, from a view controller life cycle, you should **keep a reference to the task and cancel it on deinit**

make sure to use weak self, and never capture again a strong reference inside the task, this will cause self to be strongly captured. You can use self? in the whole code inside the task to avoid memory cycle.

```
class AViewController: UIViewController {

    /// keep a reference if the task is super long
    /// if the task finished quickly, is not a big problem
    /// it will cancel itself when done
    /// AsyncStream -> ALWAYS cancel them
    private var loadingTask: Task<Void, Never>?
    /// just a worker to exemplify an async stream creation
    private var aWorker = AWorker()
    /// keep a reference to the task used to call the async stream... so you can cancel it
    /// when the view controller deinit - make sure not to make any reference cycle inside that task
    private var aTaskReference: Task<Void, Never>?

    deinit {
        loadingTask?.cancel()
        aTaskReference?.cancel()
    }

    override func viewDidLoad() {
        super.viewDidLoad()

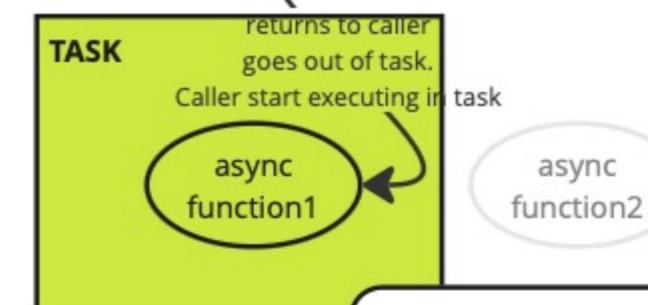
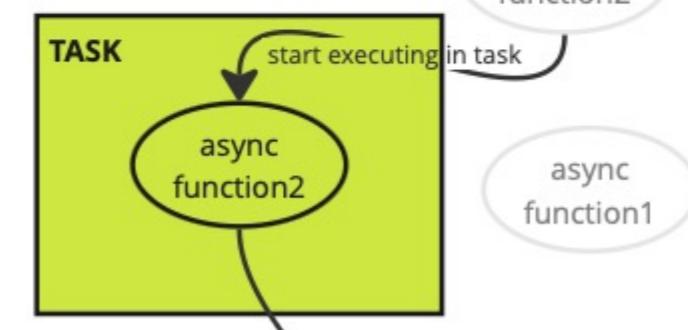
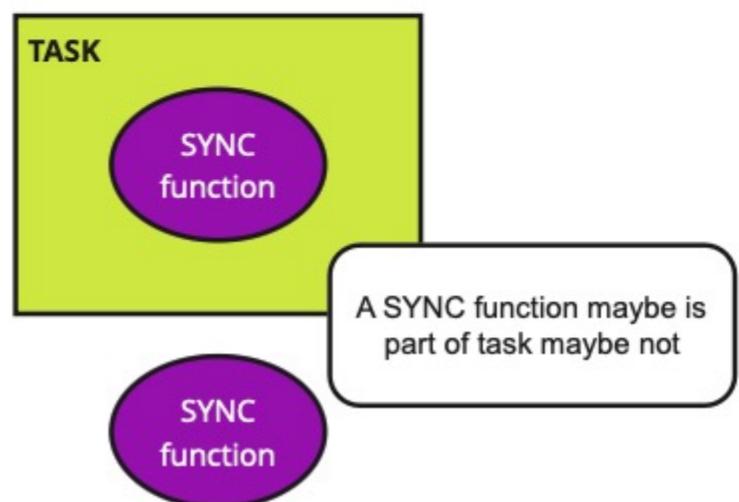
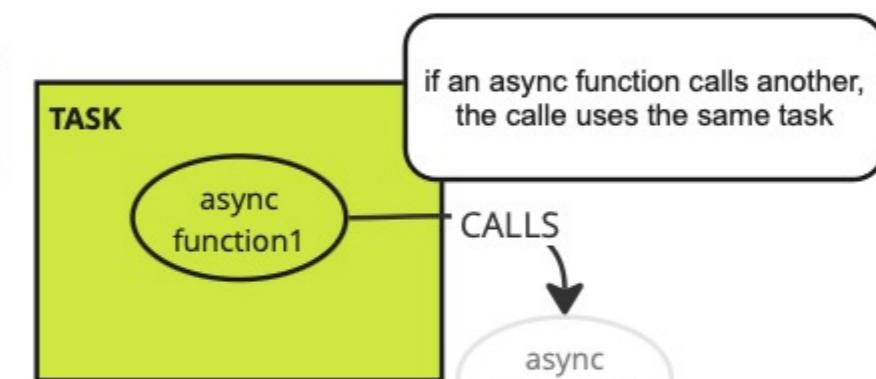
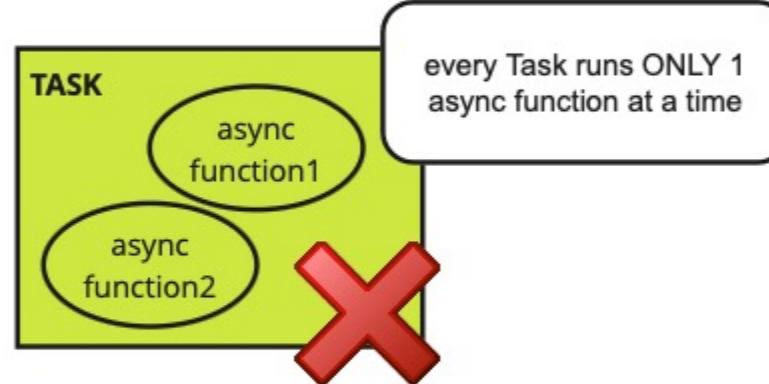
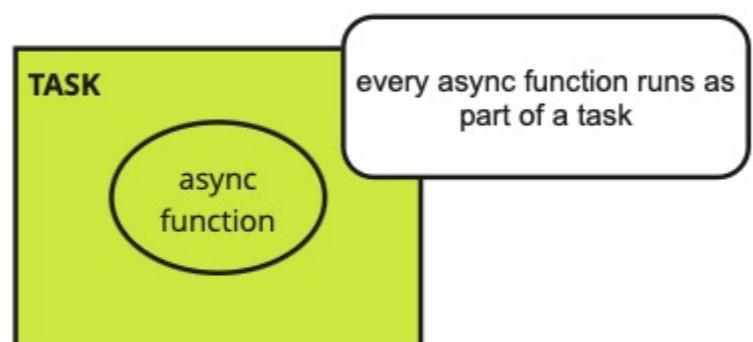
        // example for a normal Task (without AsyncStream inside)
        loadingTask = Task {
            // some very long calculation
        }
        aTaskReference = Task {
            for await value in aWorker.getAsyncStream() {
                // do something with the value
            }
        }
    }
}
```



# Structure concurrency, tasks, groups SE-0304

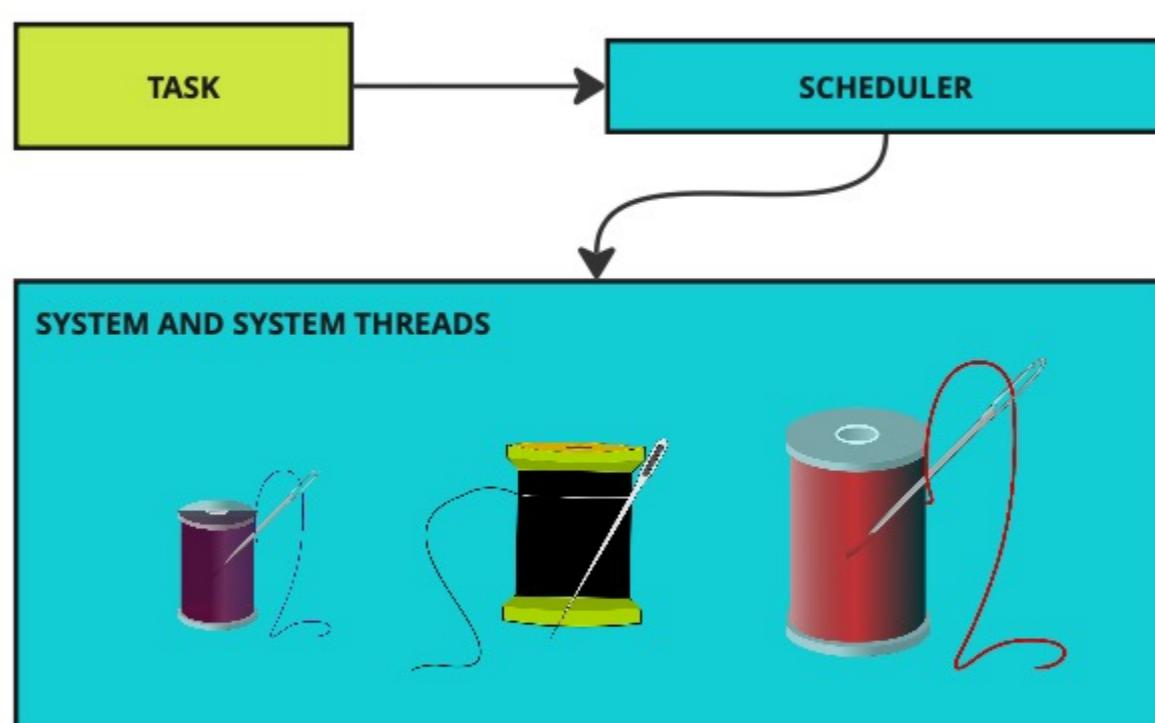
SE-0311 SE-0323 SE-0343 SE-0381 SE-417 SE-442

Structured Concurrency is the strategy that Swift offers to execute work in parallel by using Tasks and child Task. And hiding low-level thread management

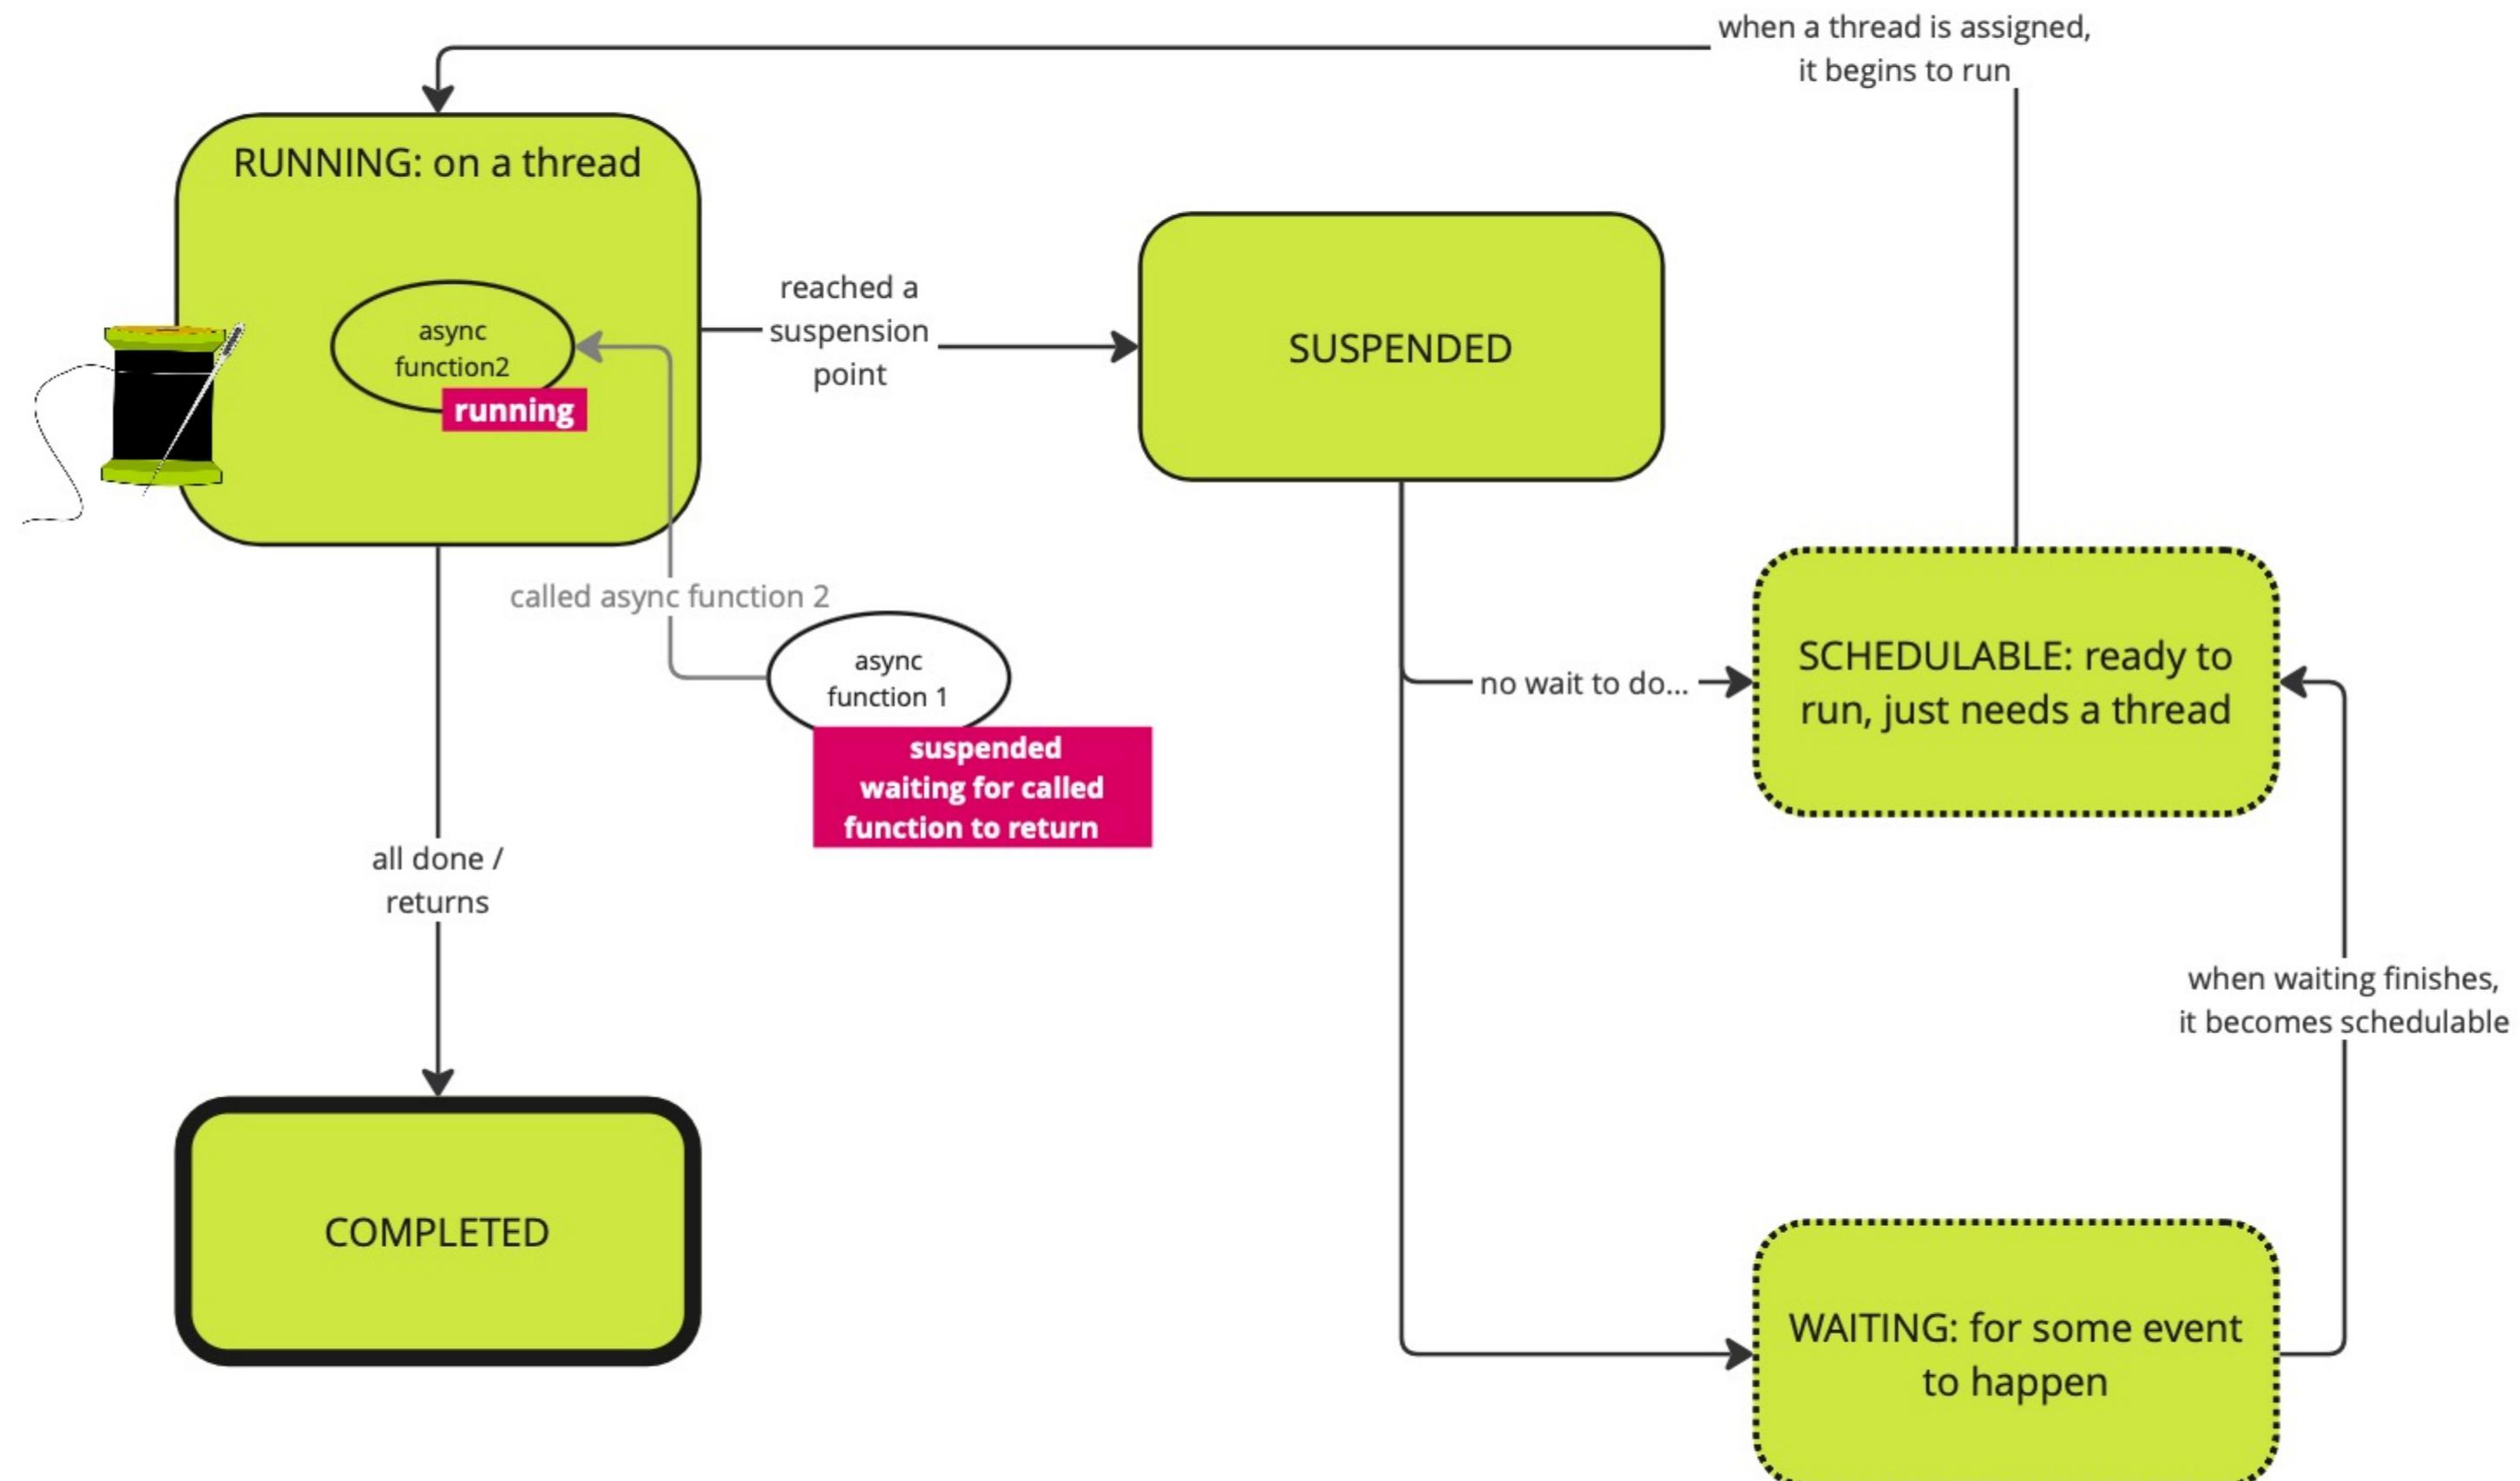


when the called async functions finished, the function 1 is again executing in Task

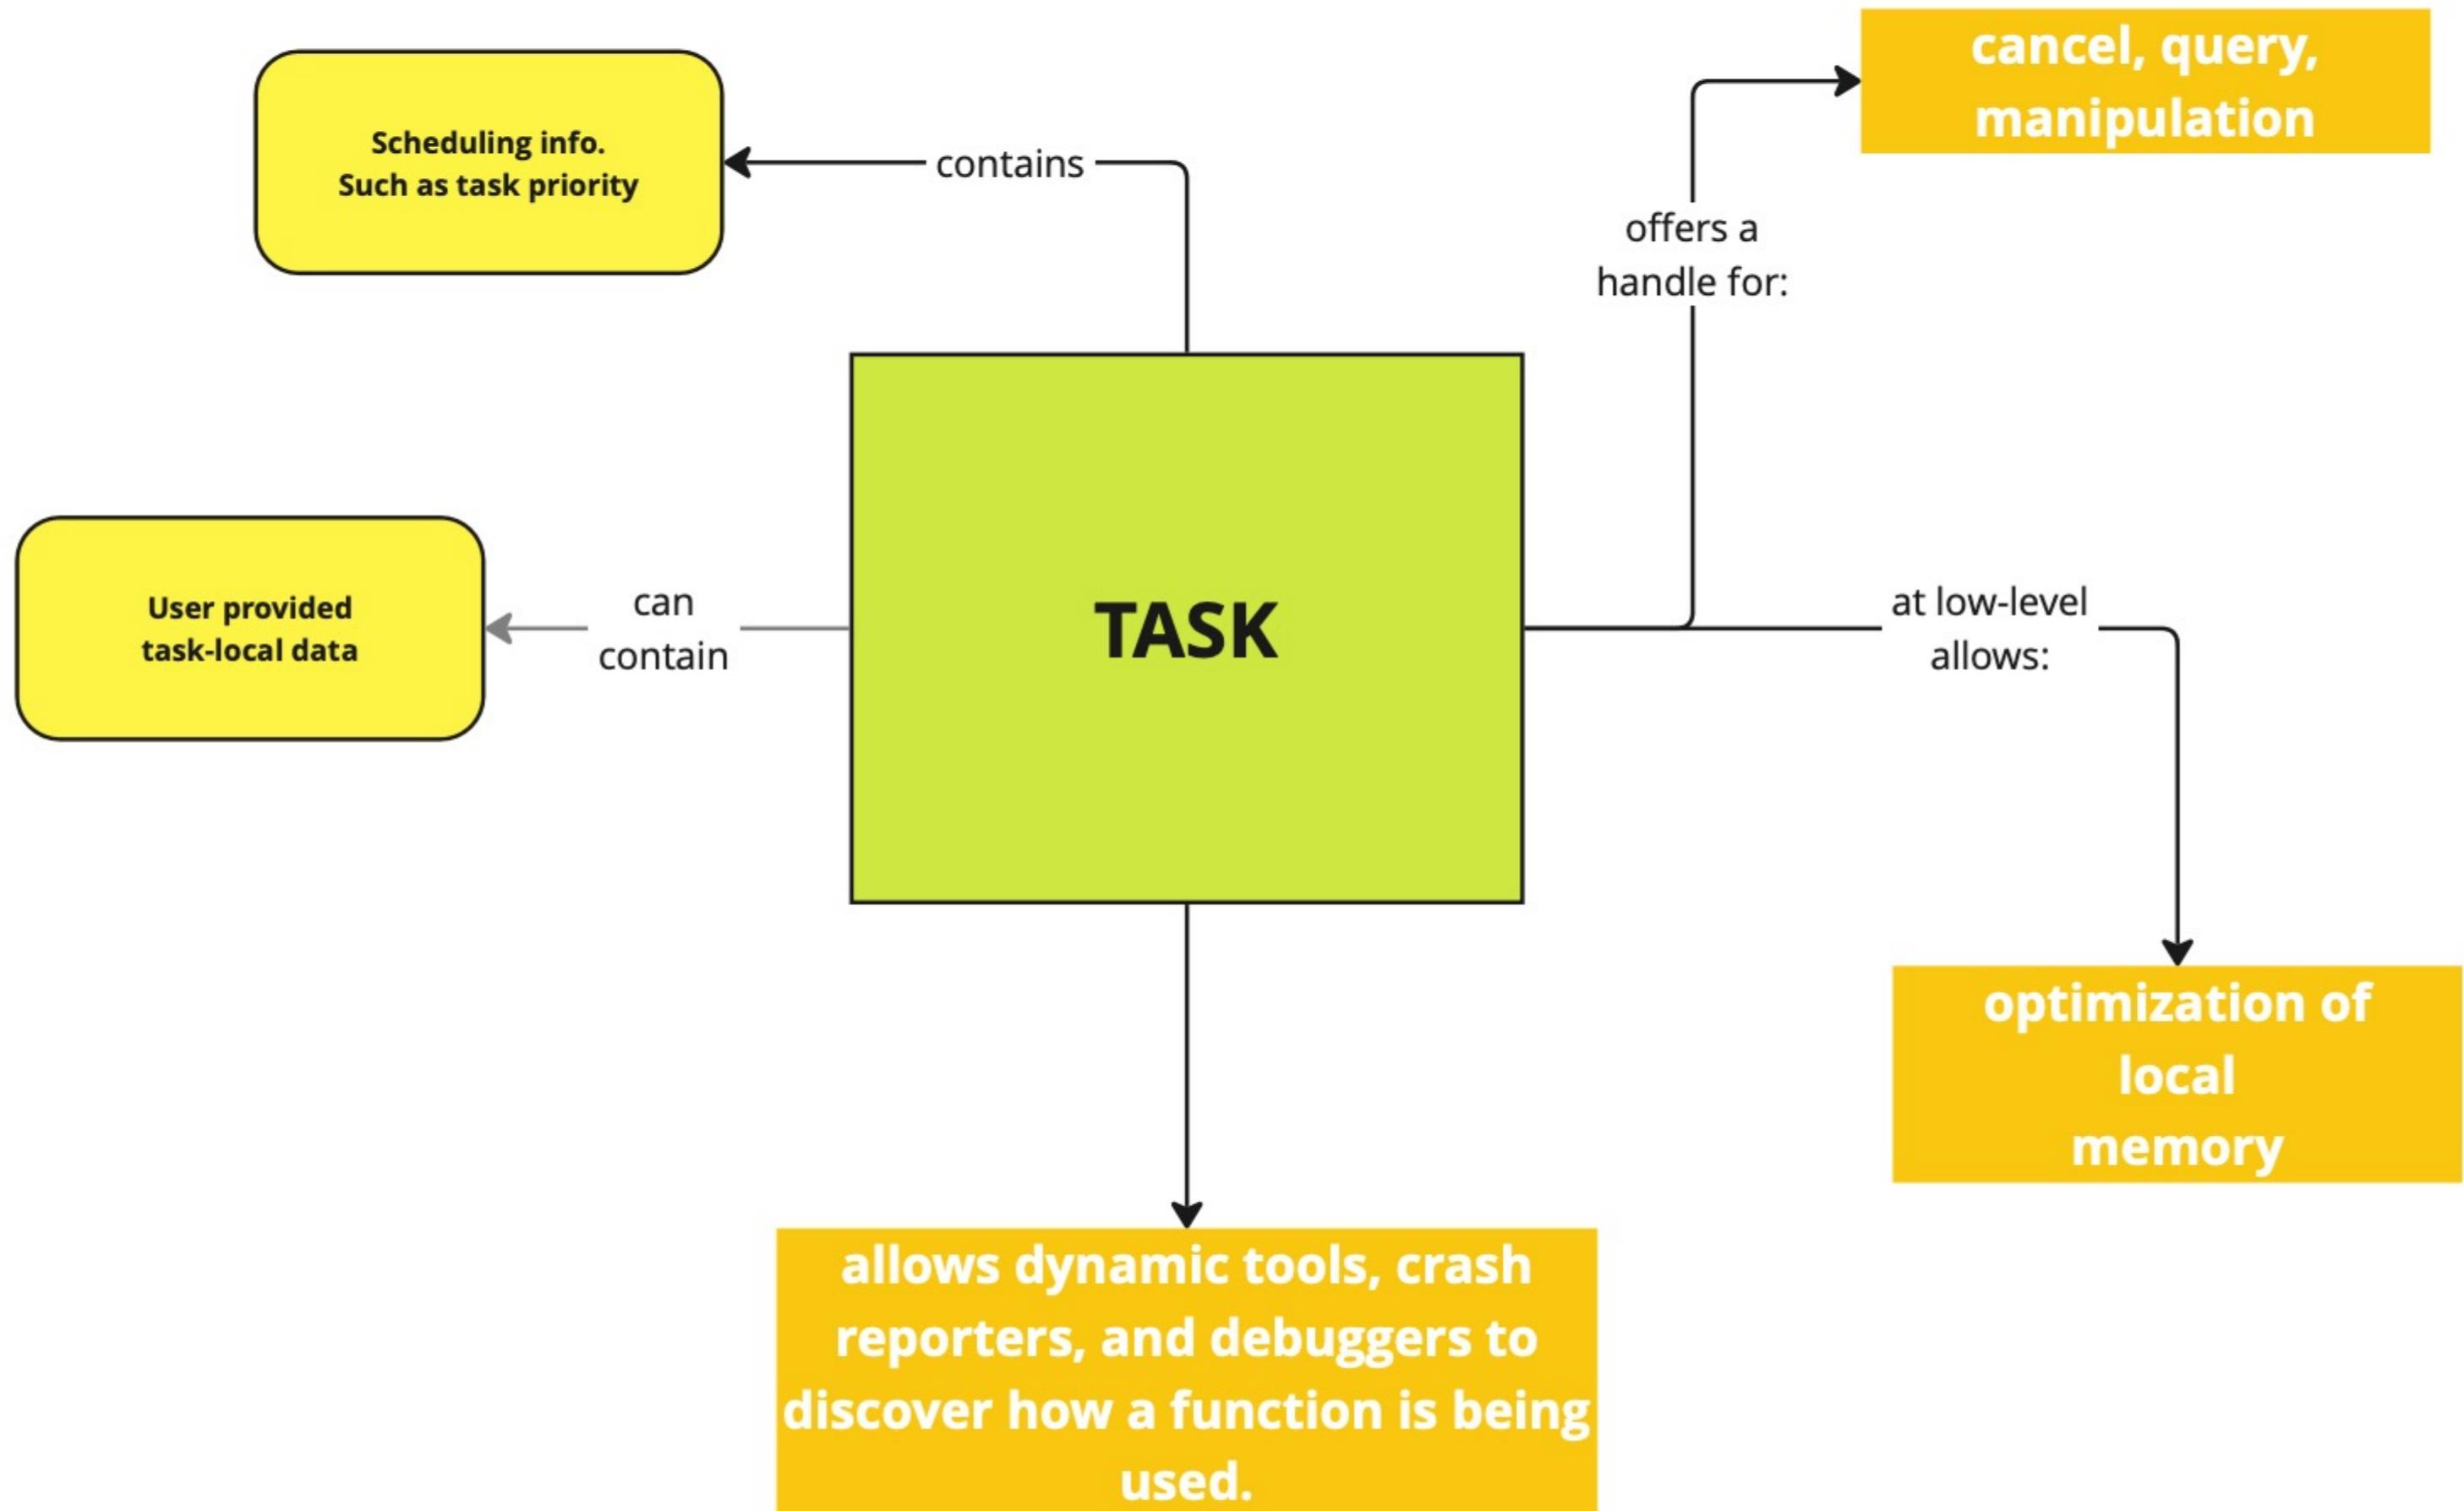
The scheduler is responsible for running a task in the system and assigning threads



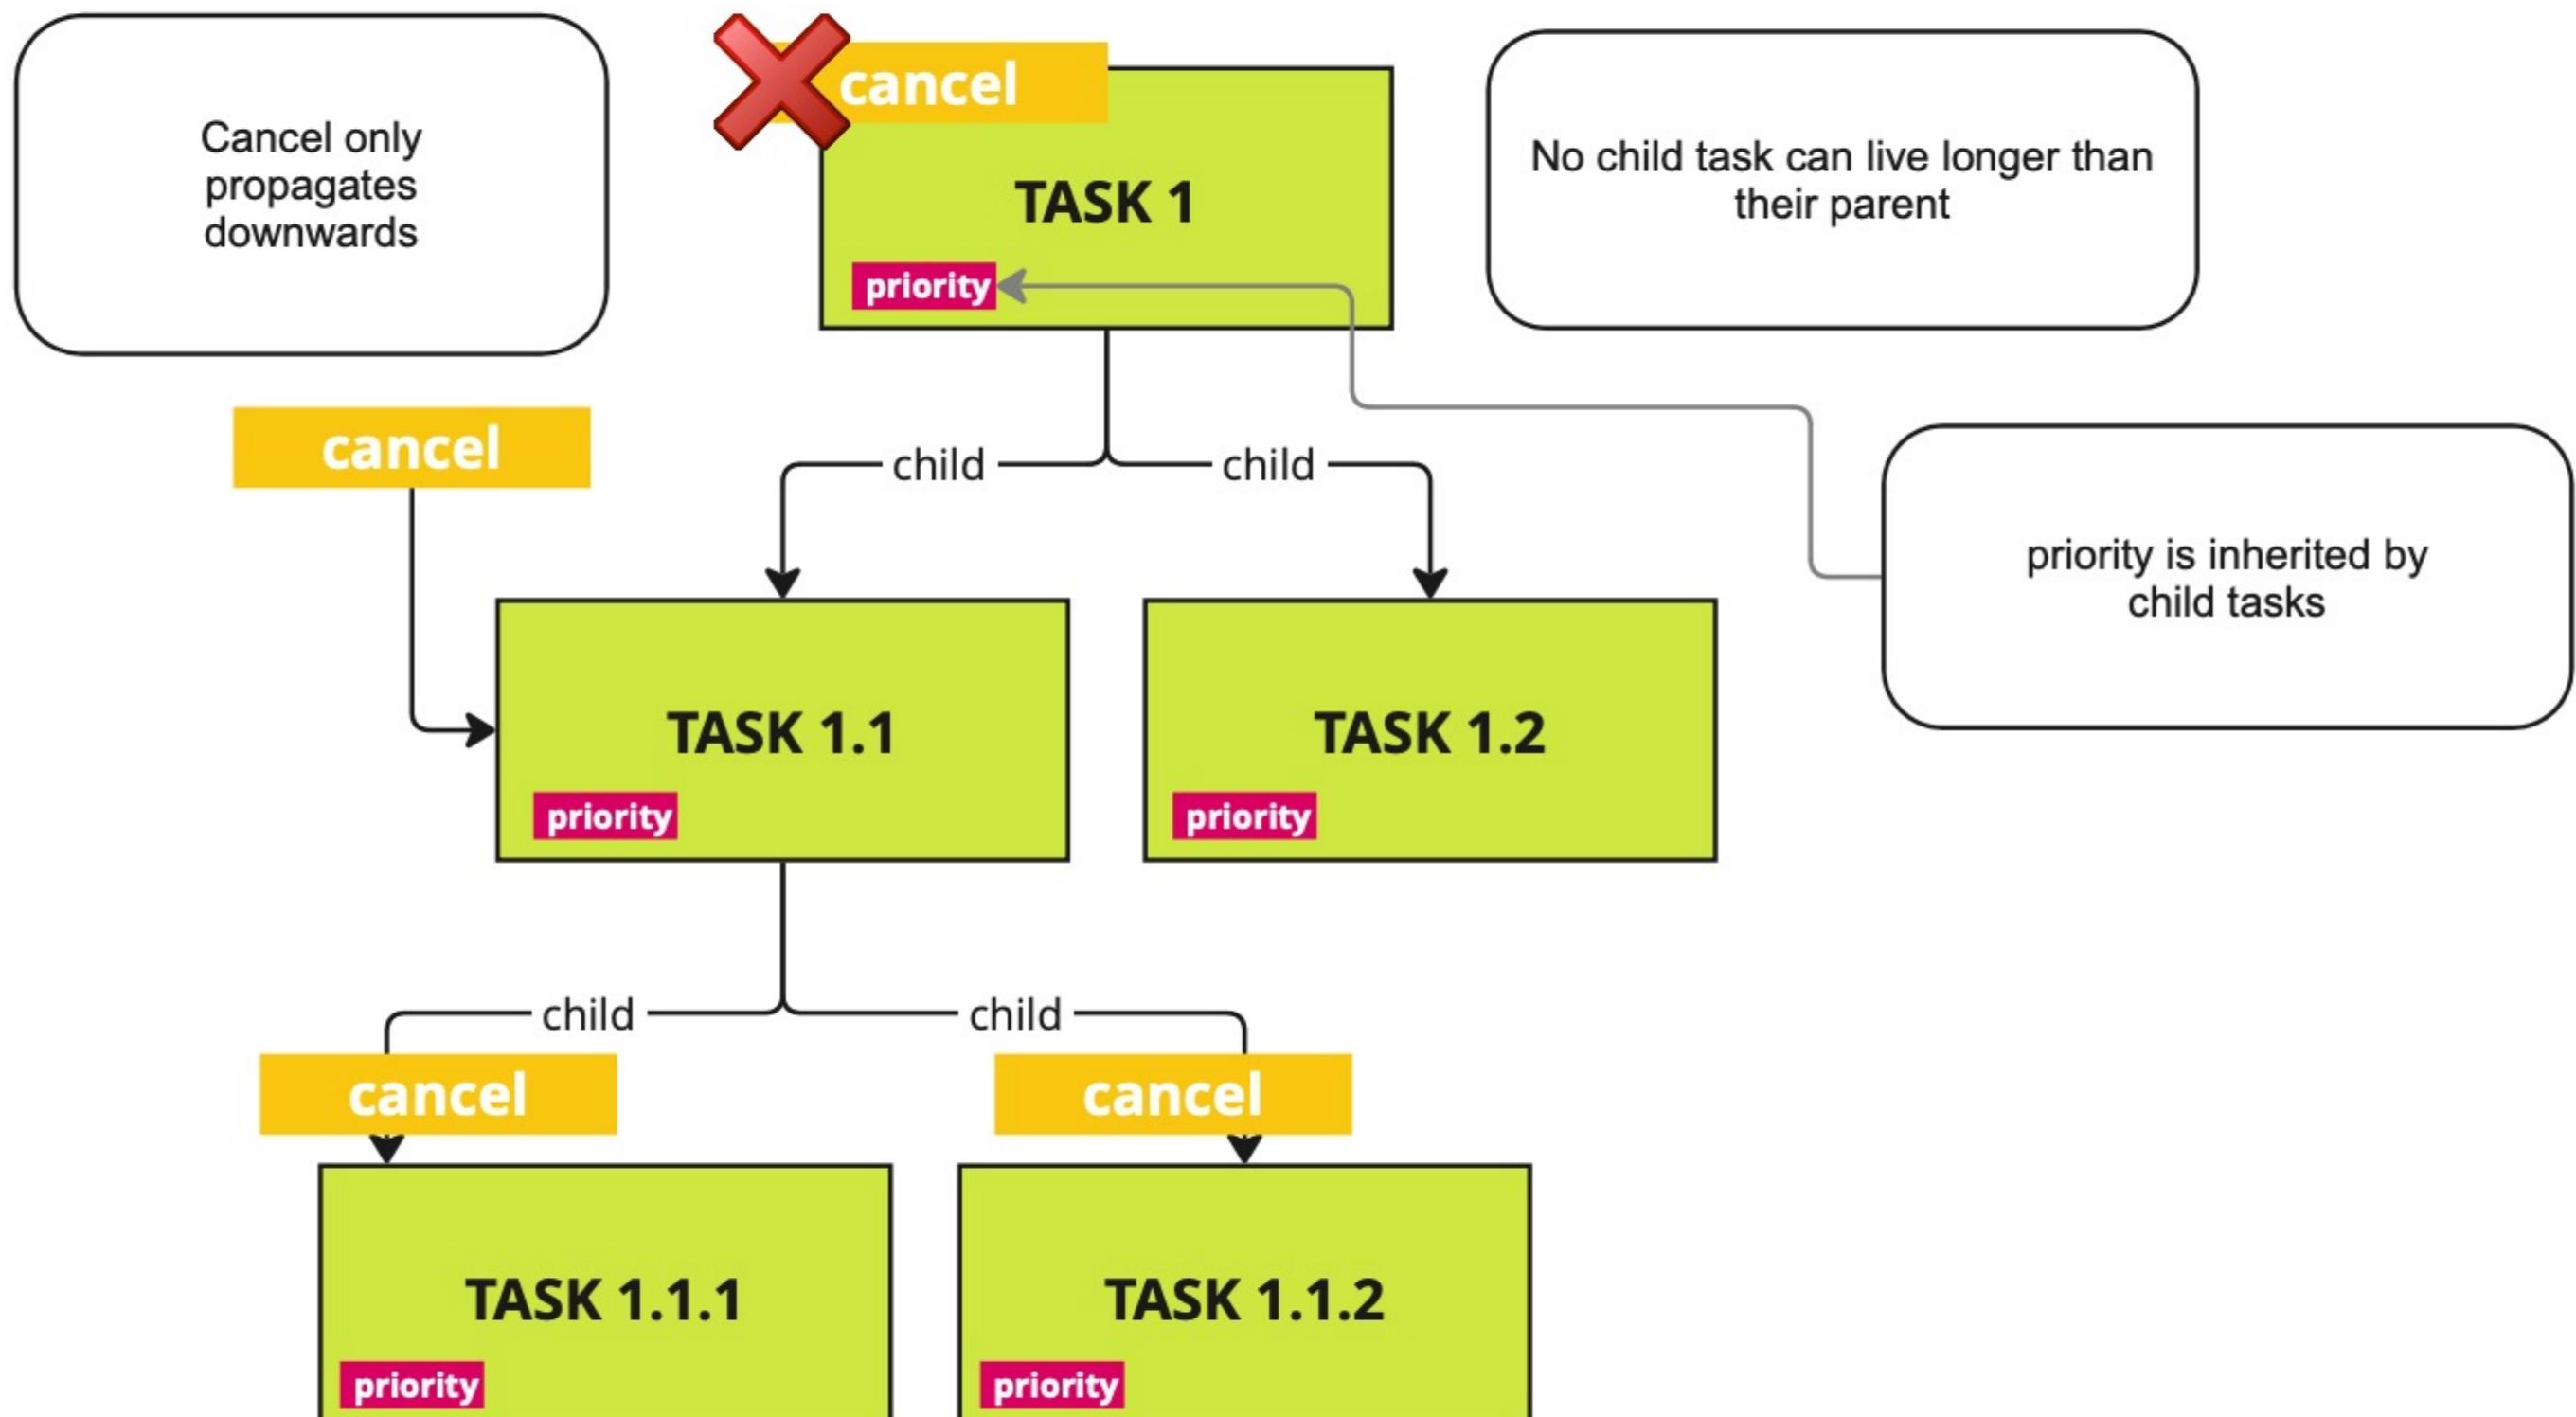
# Task's States



# Task's Structure

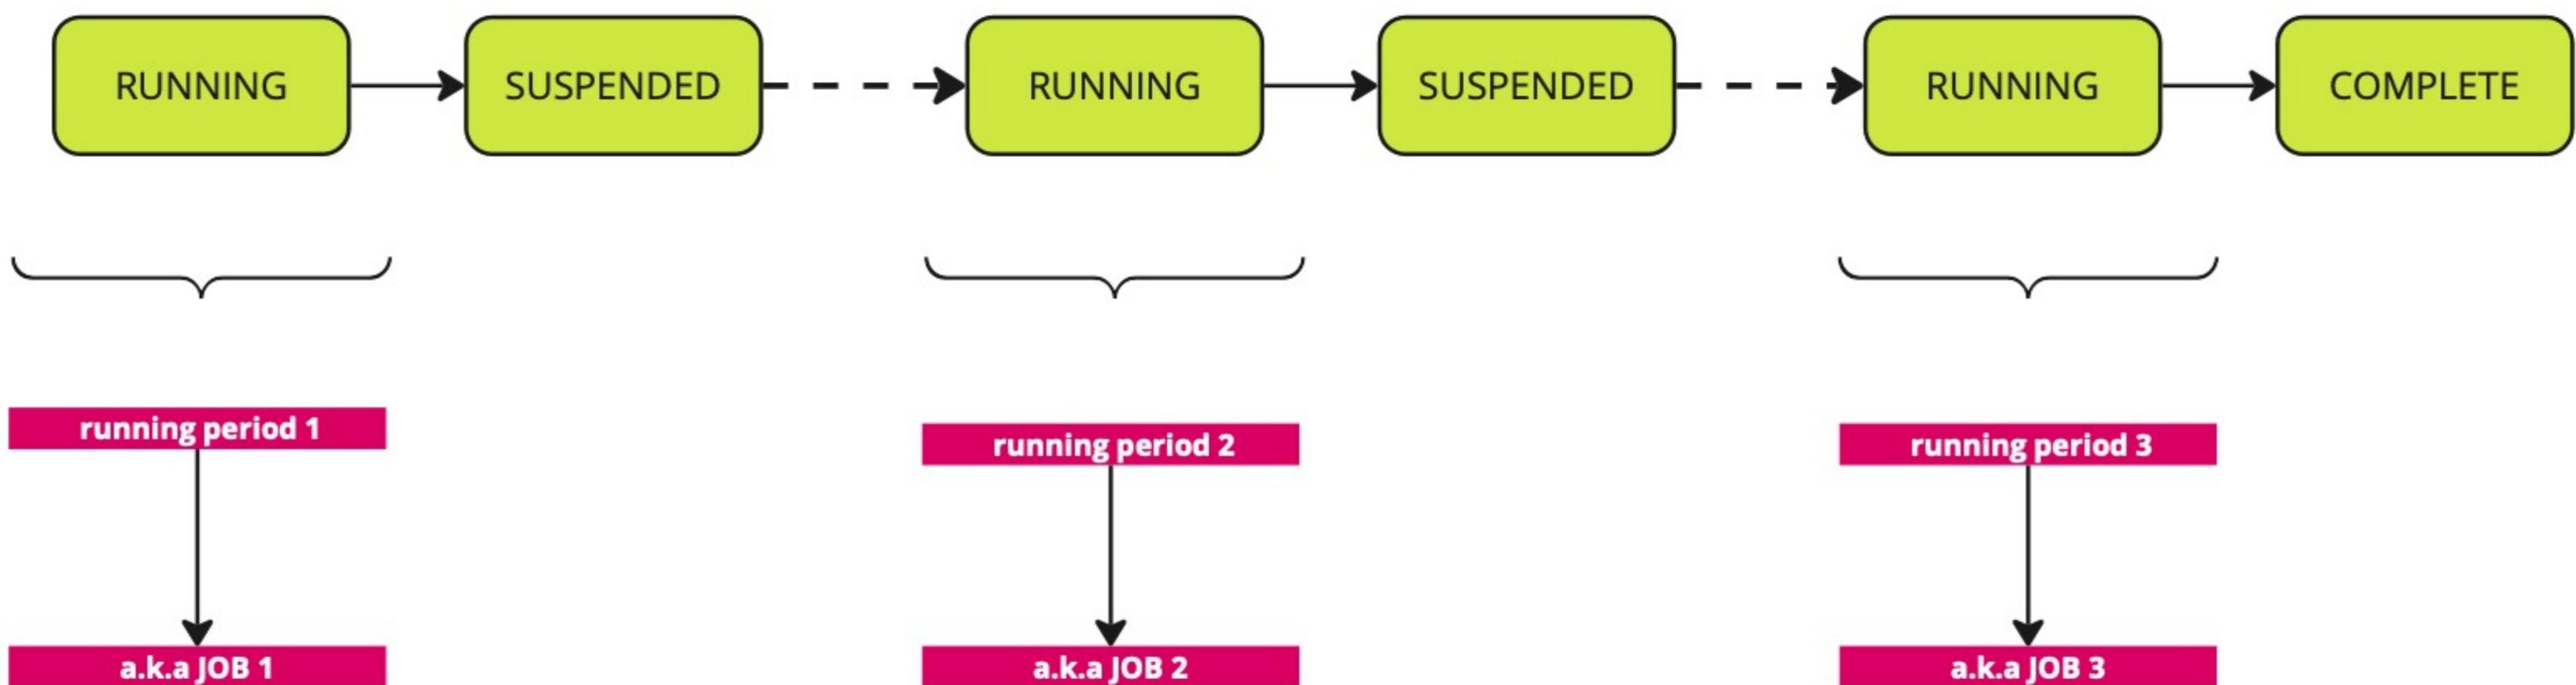


# Tasks and child tasks



# Taks and jobs

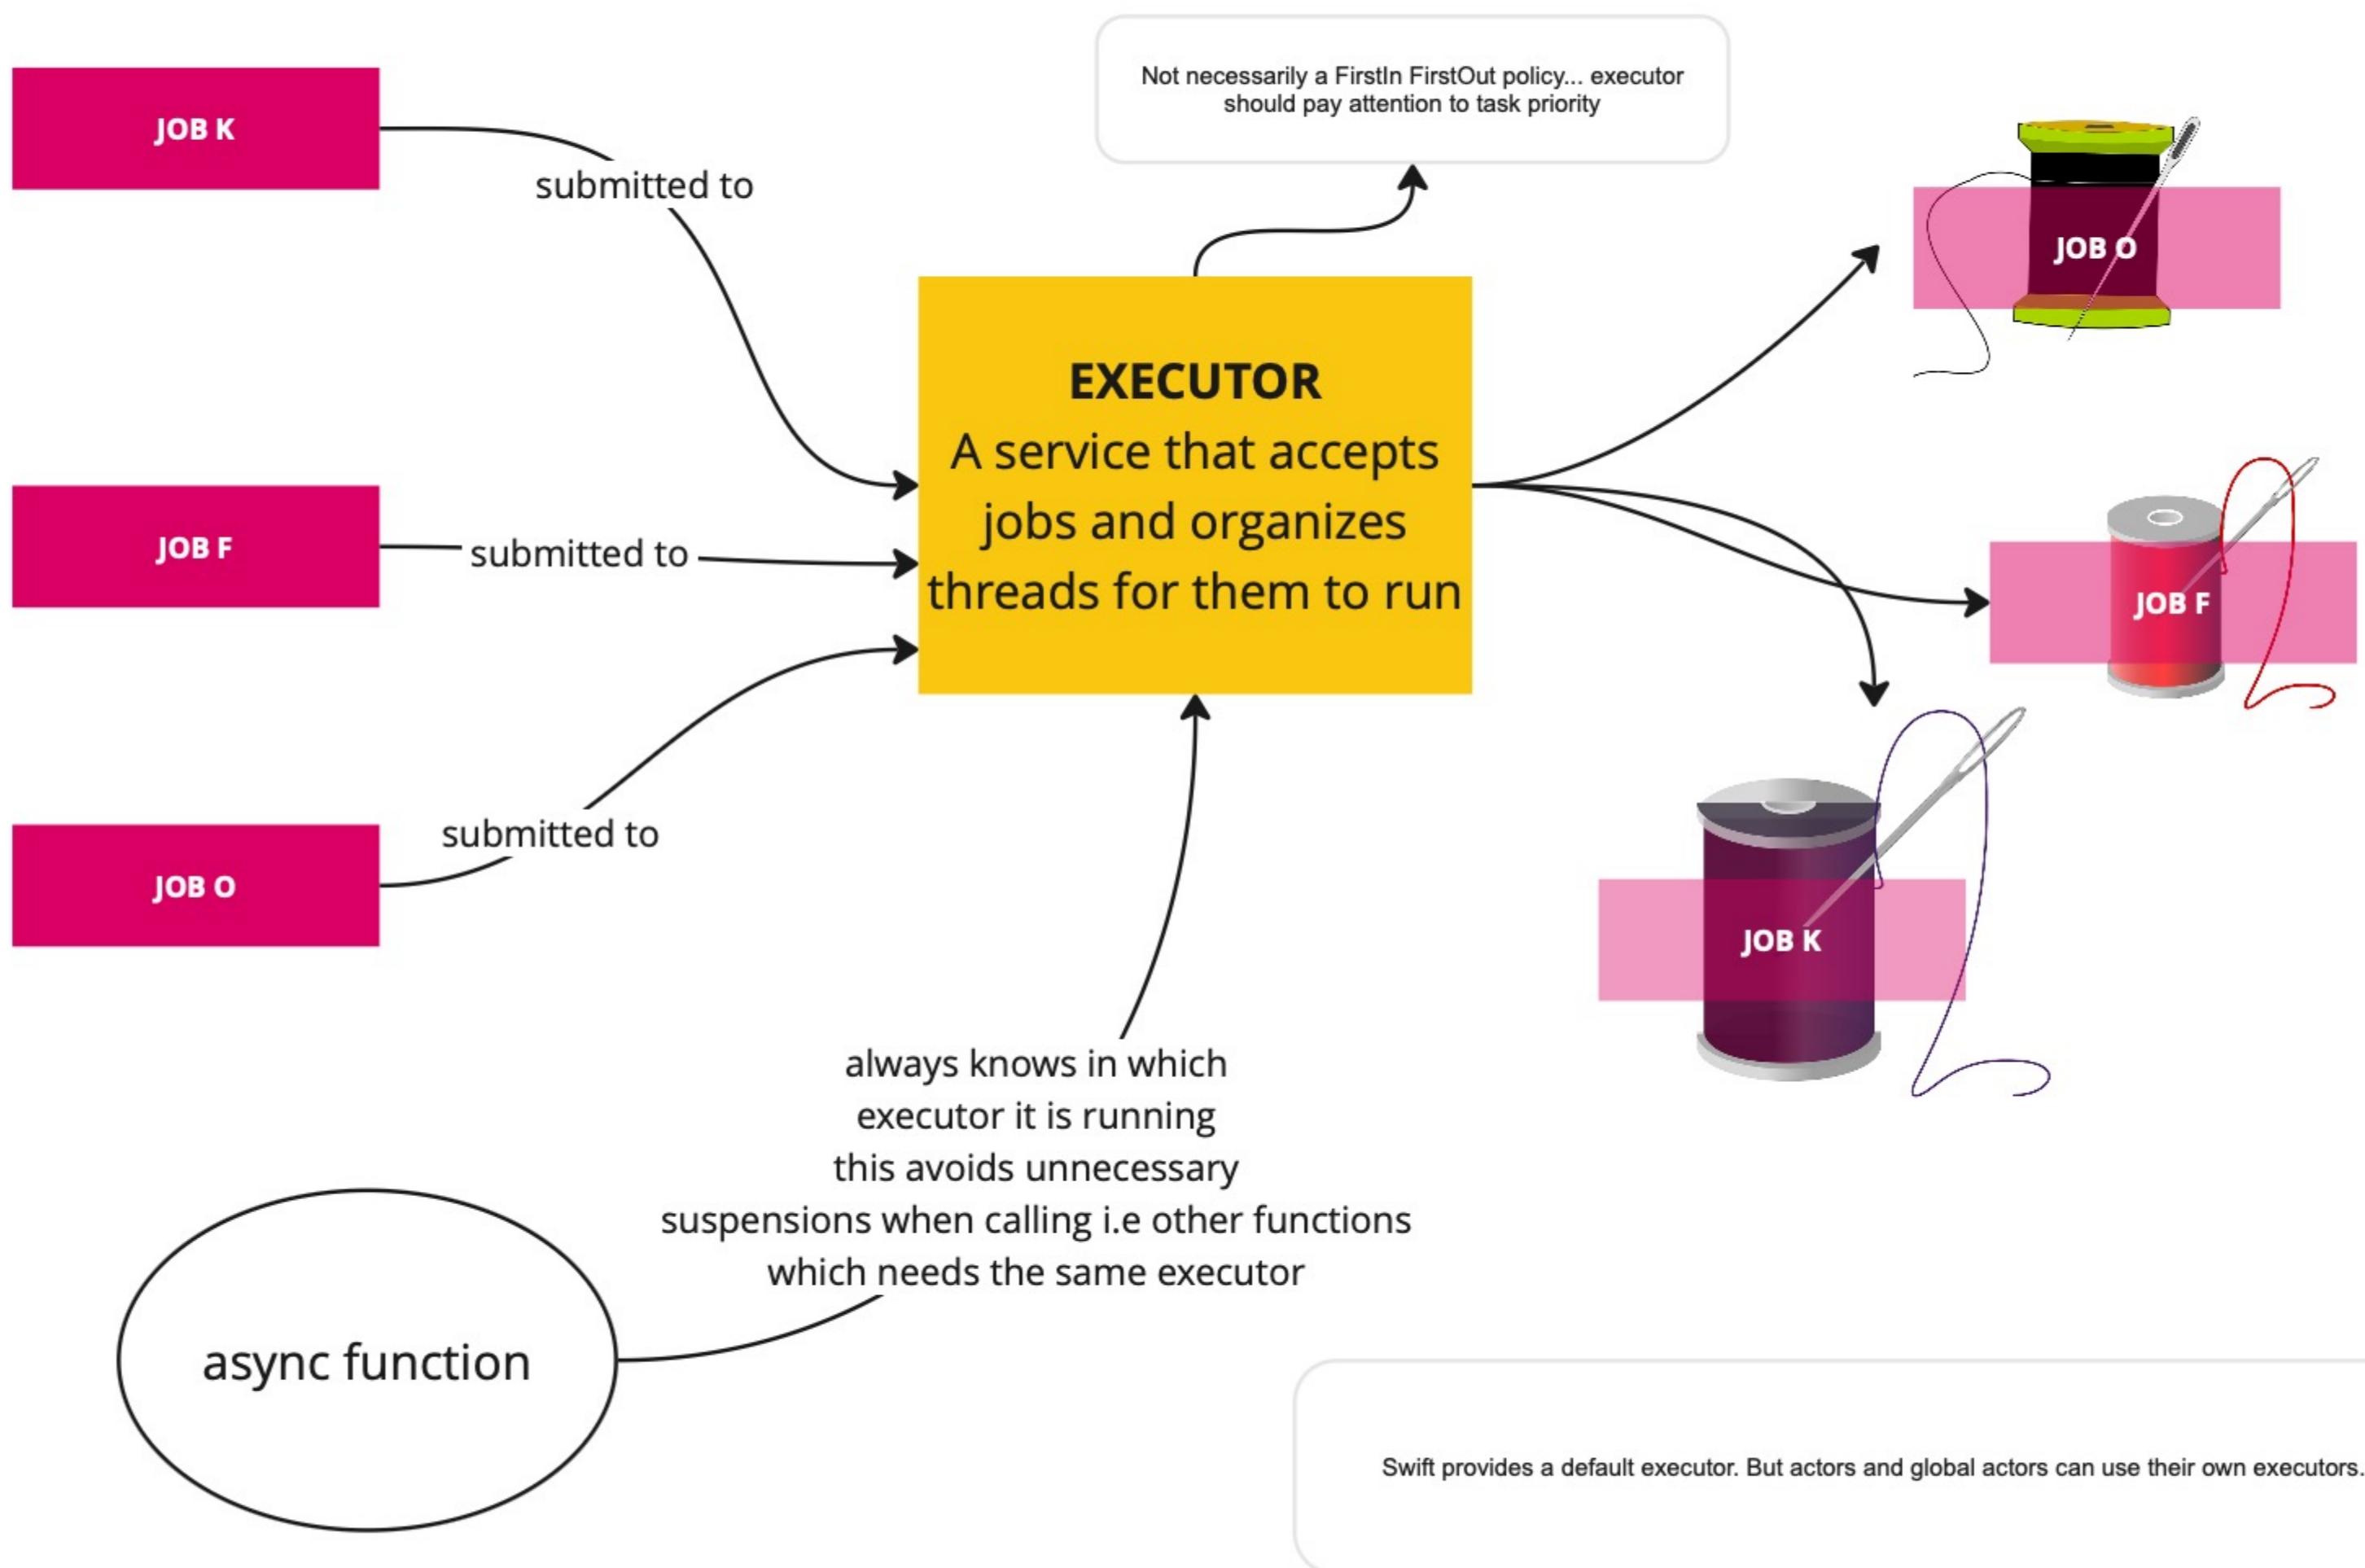
A task run in a succession of periods, till it is suspended, and then runs again.  
Those running periods are called JOBS



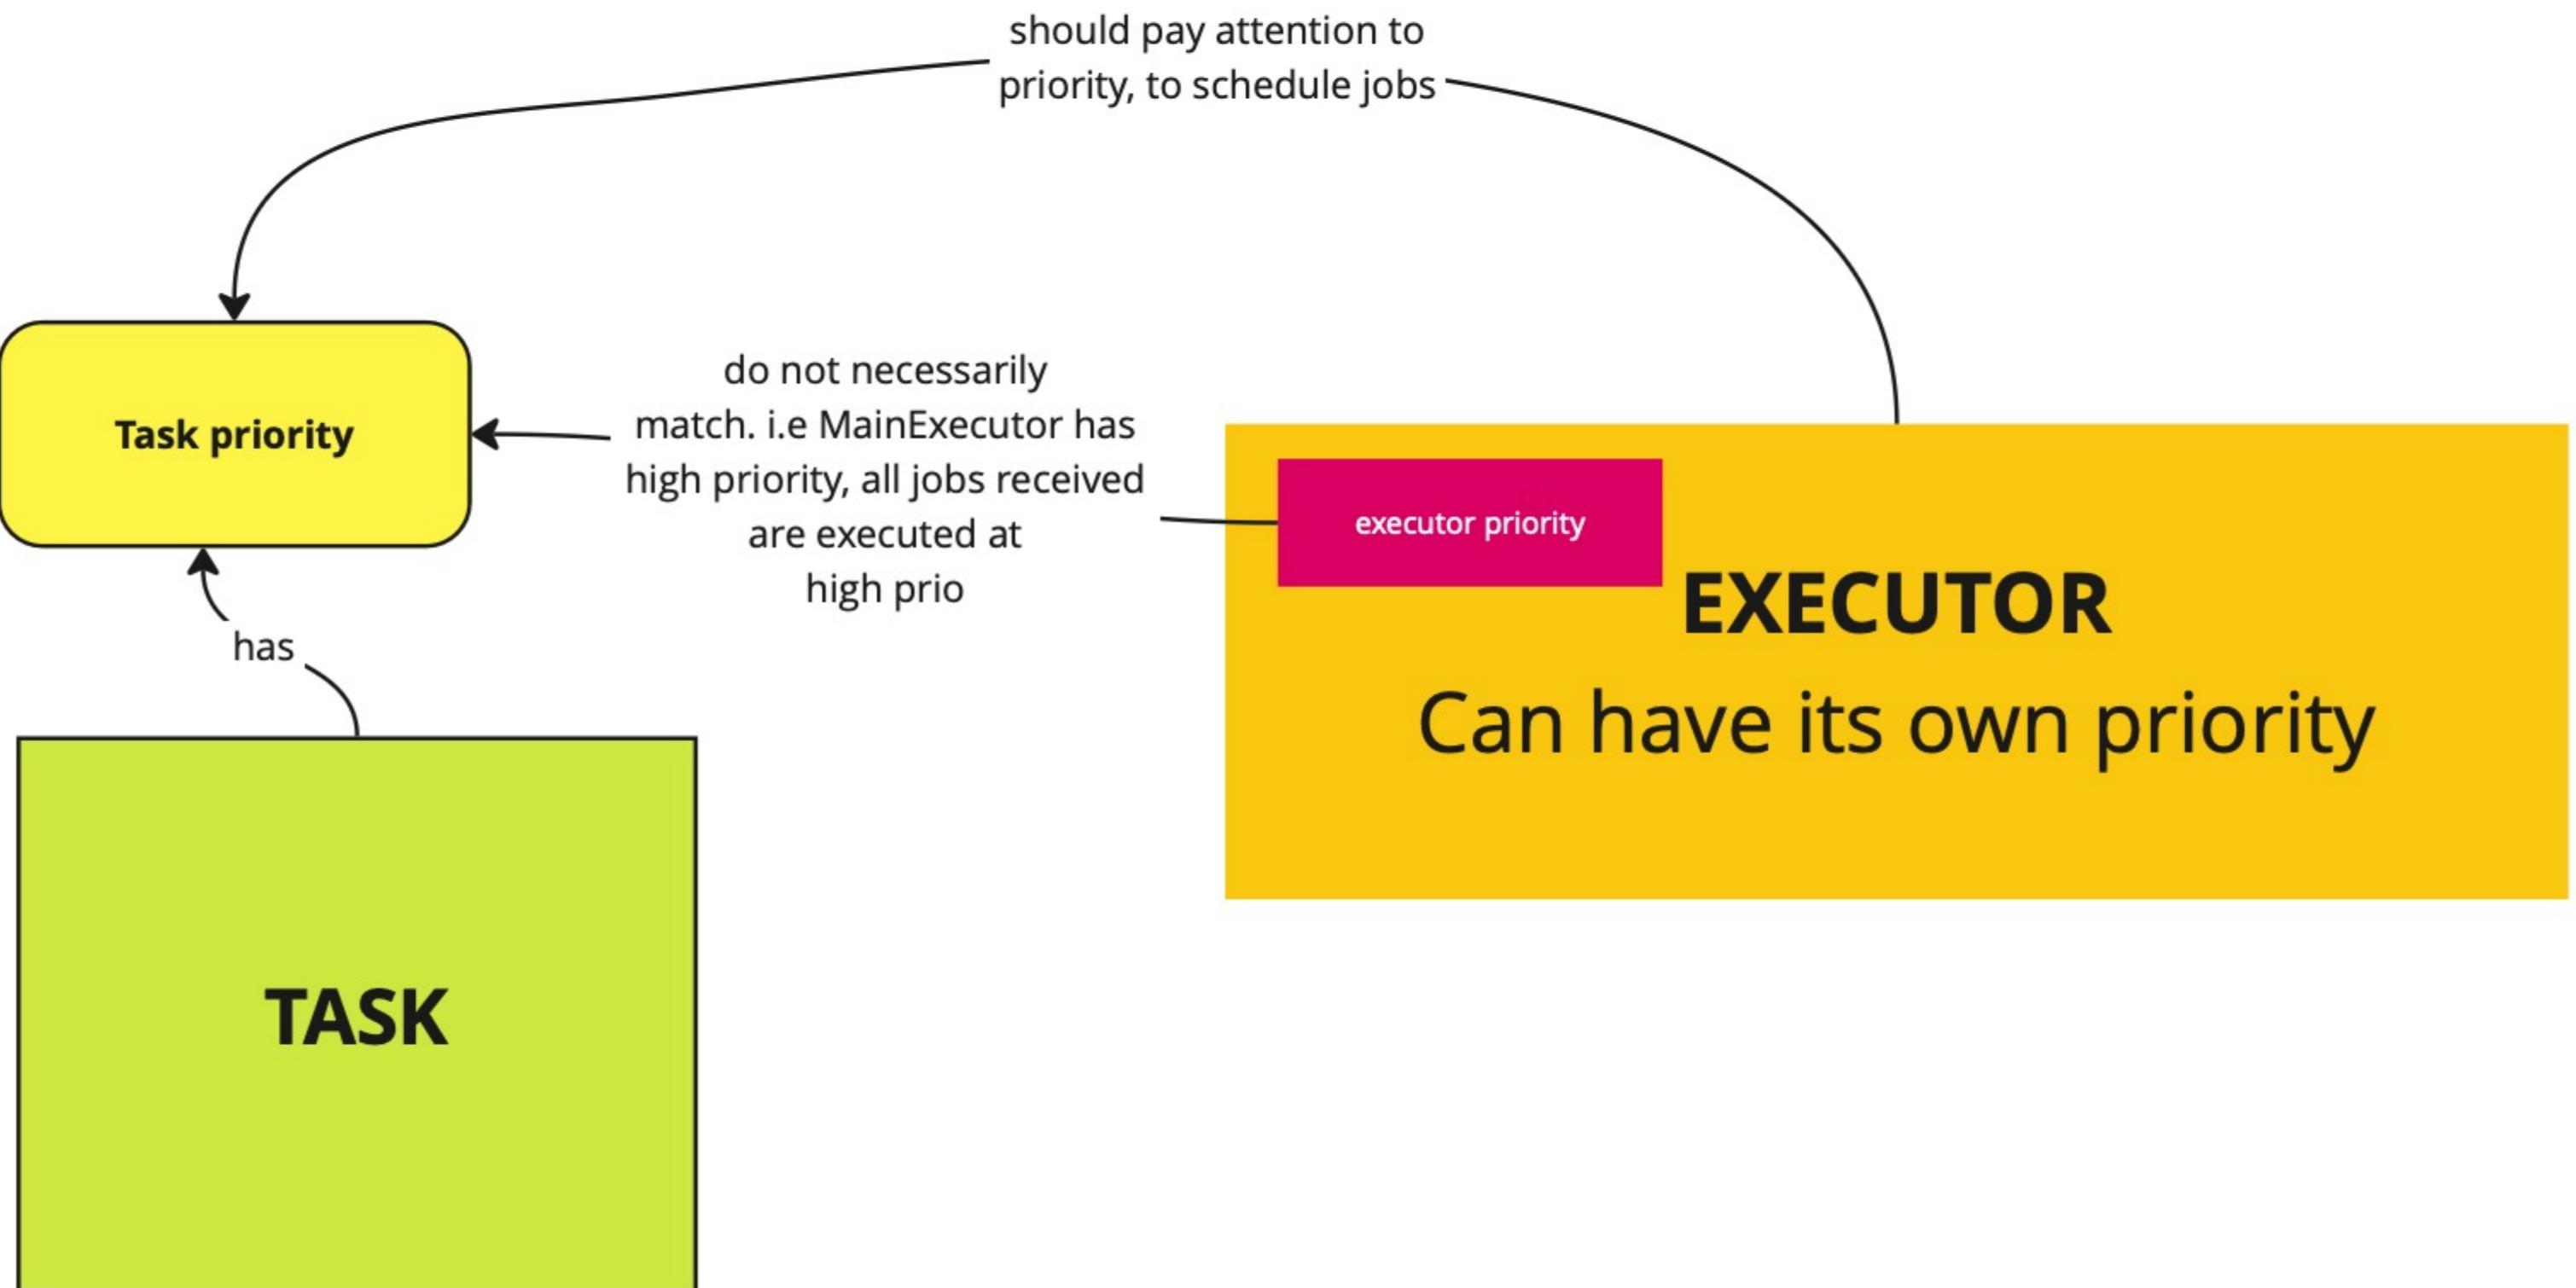
**JOB = BASIC UNIT OF SCHEDULABLE WORK**

YOU ONLY NEED TO BE AWARE OF JOBS WHEN YOU IMPLEMENT YOUR OWN EXECUTOR

# Executors



# Executors and priority



## Task and main

```
// Swift will create a new task that will execute main().  
// Once that task completes, the program terminates.  
@main  
struct SwiftConcurrencyConceptualMapsApp: App {  
    var body: some Scene {  
        WindowGroup {  
            ContentView()  
        }  
    }  
}
```

# Task and cancellation

CAN BE  
**EXPLICIT**

BY CALLING CANCEL on a task  
handler

CAN BE  
**AUTOMATIC**

when the parent task is cancelled

**TASK**

**flag = cancelled**

on cancellation a flag is set forever on the task, that says it is cancelled

```
struct Vegetable: Ingredient {  
    func chopped() async throws -> Vegetable {  
        try? await Task.sleep(for: .seconds(0.3))  
        // if this task would be long, we need to check for cancellations  
        try Task.checkCancellation()  
        // if it is still not cancelled, we can continue  
        return self  
    }  
}
```

if you register a cancellation handler, then it will be called  
immediately after cancellation, even if inside the code you do not  
check for cancellation

```
loadingTask = Task {  
    // some very long calculation here....  
    // or if you need to react immediately to  
    // the cancellation, create a  
    // task with cancellation handler  
    await withTaskCancellationHandler {  
        // some long processing here  
    } onCancel: {  
        // do some cleanup  
    }  
}
```

# Unstructured task

primary rule of structured concurrency: that a *child task* cannot live longer than the *parent task*

**therefore**

That's why we need Unstructured tasks. For fire-and-forget code. Or to call async code from a sync code



can be...



**non-detached**

they inherit context, depending on where they are called from



**Called from a sync function which is inside a task:** inherit the priorities, task local values (a copy of them),

**Called from an actor function:** all the same as when called from a sync function inside a task + the actor executor context + access to the actor isolated state

**detached: does not inherit anything**



**Called from code not running in any Task:** infer priority from the system + is not isolated to anything + executes in the concurrent global actor

## About TaskGroup and DiscardingTaskGroup

Task groups are the building block of structured concurrency. They enable features such as automatic cancellation propagation, propagating errors, and ensuring well-defined lifetimes

**but they have a limitation**  
since they are expected to return the results, they should keep them in memory until they are consumed, this can blow out the memory space if results are not consumed and a task group is potentially infinite (i.e. server attending requests loop)

some attempts to solve this, before the addition of DiscardingTaskGroups involved limiting the max capacity of a group... but manually

```
try await withThrowingTaskGroup(of: Void.self) { group in
    // Fill the task group up to maxConcurrency
    for _ in 0..<maxConcurrency {
        guard let newConnection = try await listeningSocket.accept() else {
            break
        }

        group.addTask { handleConnection(newConnection) }
    }

    // Now follow a one-in-one-out pattern
    while true {
        _ = try await group.next()
        guard let newConnection = try await listeningSocket.accept() else {
            break
        }
        group.addTask { handleConnection(newConnection) }
    }
}
```

Discarding groups, do this automatically... they clean up tasks when completed, BUT therefore can not have a next() function NOR have ChildTaskResult

```
// GOOD, no leaks!
try await withThrowingDiscardingTaskGroup() { group in
    while let newConnection = try await listeningSocket.accept() {
        group.addTask {
            handleConnection(newConnection)
        }
    }
}
```

# Where do async functions run?

SE-0296

now we have  
**async functions**

SE-0306

now we have  
**actors and**  
**actor-isolated async**  
**functions may run in their**  
**actor's executor**  
and also non-isolated  
functions won't change  
executor, so they may also  
use the actor's executor

SE-0338

but then **non-isolated**  
**async functions** may over-  
populate the actor's  
executor...  
so now let's make them  
**FORMALLY run in a generic,**  
**non-actor global executor**

HOP-OFF  
SEMANTIC  
INTRODUCTION

formally = in theory... in  
reality internally the  
system can decide if to do  
the switch or not for i.e  
optimization

SE-0417

Swift 6.0

but now, there will be  
lots of HOP-OFFS  
potentially...so if you need  
to not switch context so  
often... **TASK EXECUTOR**  
**PREFERENCE** is introduced  
and you can choose

SE-0420

Swift 6.0

but these hop-offs can be  
very inefficient... so let's give  
developers the chance to do  
something to indicate that  
the actor context has to be  
inherited also for non-  
isolated functions....**ACTOR**  
**EXECUTOR INHERITANCE**